

Tongzhou Wang

✉ tongzhou@mit.edu | 🏠 tongzhouwang.info | 📄 Google Scholar | 📷 ssnl

EDUCATION

Massachusetts Institute of Technology

Ph.D. in Computer Science

2019 - 2024

- Advisors: Antonio Torralba, Phillip Isola
- Thesis: Intelligent Agents via Representation Learning

Massachusetts Institute of Technology

M.S. in Computer Science

2019 - 2022

- Advisors: Antonio Torralba, Phillip Isola
- Thesis: Geometric Properties of Learned Representations

University of California, Berkeley

B.A. in Computer Science and Statistics

2013 - 2017

- Advisors: Stuart J. Russell, Ren Ng, Alexei A. Efros

EXPERIENCES

Simons Institute for the Theory of Computing

Long-Term Visitor

2024

- Summer Cluster: AI, Psychology, and Neuroscience.

Facebook AI Research (FAIR), Meta Platforms, Inc.

Part-Time Student Researcher

2021

- Mentor: Yuandong Tian. Minimal world model for reinforcement learning. Paper published in ICML 2022.

Facebook AI Research (FAIR), Meta Platforms, Inc.

Research Intern (FAIR)

2021

- Mentor: Yuandong Tian. Research on world model for reinforcement learning.

Facebook AI Research (FAIR), Meta Platforms, Inc.

Software Engineer

2017 - 2019

- Built PyTorch, a leading software framework for deep learning. Works on data pipelines, autograd, ML operators, etc.

RESEARCH INTERESTS

Learning world representations for generalist agents.

I am interested in learning structured representations that aggregate and select information about the world from various data sources, improve multi-task training, and enable autonomous adaptation to new tasks.

FEATURED PUBLICATIONS

(* indicates equal contribution)

The Platonic Representation Hypothesis

Minyoung Huh*, Brian Cheung*, Tongzhou Wang*, Phillip Isola*

2024

International Conference on Machine Learning 2024 [[ICML 2024 \(Position Paper\)](#)]

Optimal Goal-Reaching Reinforcement Learning via Quasimetric Learning

Tongzhou Wang, Antonio Torralba, Phillip Isola, Amy Zhang

2023

International Conference on Machine Learning 2023 [[ICML 2023](#)]

Denoised MDPs: Learning World Models Better Than the World Itself

Tongzhou Wang, Simon S. Du, Antonio Torralba, Phillip Isola, Amy Zhang, Yuandong Tian

2022

International Conference on Machine Learning 2022 [[ICML 2022](#)]

Understanding Contrastive Representation Learning through Alignment and Uniformity on the Hypersphere

Tongzhou Wang, Phillip Isola

2020

International Conference on Machine Learning 2020 [[ICML 2020](#)].

Dataset Distillation

Tongzhou Wang, Jun-Yan Zhu, Antonio Torralba, Alexei A. Efros

2018

Learning to See by Looking at Noise

Manel Baradad*, Jonas Wulff*, Tongzhou Wang, Phillip Isola, Antonio Torralba

2021

Advances in Neural Information Processing Systems 2021 [[NeurIPS 2021](#)]

INVITED TALKS

Reinforcement Learning as Representation Learning

Simons Institute
UC Berkeley

June 2024
May 2024

Structured Representations for Active Agents

Stanford Vision and Learning Lab (SVL), Stanford University
Guest Lecture, USC

November 2023
November 2023

Quasimetric Reinforcement Learning

Brown University
AI Seminar, CMU
Vector Institute
Deep Learning: Classics and Trends (DLCT)
Machine Learning Advances Symposium, MIT
UT Austin
Northeastern University

November 2023
October 2023
September 2023
June 2023
May 2023
April 2023
April 2023

Technical Talks on PyTorch Internals

PyTorch Developer Conference, San Francisco, CA, USA
Global Mobile Internet Conference, Beijing, China

October 2019
April 2018

MENTORING

Massachusetts Institute of Technology

Adrian Rodriguez Munoz (Ph.D. student)
Hyojin Bahng (Ph.D. student)
David X. Wu (B.S. & M.S. '22; now Ph.D. student at UC Berkeley)
Jingwei Ma (B.S. & M.S. '21; now Ph.D. student at University of Washington)
Steven Liu (B.S. & M.S. '21; now at TwoSigma)

Spring 2024 - PRESENT
Summer 2023 - PRESENT
Summer & Fall 2021
2019 - 2022
2019 - 2020

Carnegie Mellon University

George Cazenavette (M.S. '22; now Ph.D. student at MIT)

2021 - 2023

Summer Geometry Initiative (SGI)

Daniel Perazzo (master student at IMPA, Brazil)
Biruk Abere (B.S. student at University of Gondar, Ethiopia)
Gabriele Dominici (master student at University of Cambridge, UK)
Sana Arastehfar (master student at Queen's University, Canada)
Sanowar Raihan (research assistant at Center for Computational & Data Sciences, Bangladesh)

Summer 2023 - PRESENT
Summer 2023
Summer 2023
Summer 2023
Summer 2023

TEACHING

6.S898: Deep Learning, MIT

Teaching Assistant (Co-Designed Curriculum and Assignments for 1st Undergraduate Offering)

Fall 2022

Professional Development Course on Deep Learning, MIT

Lab Session Instructor

Summer 2019

Deep Learning Tutoring

Volunteer Tutoring for a Data Science Professional in Boston, MA, USA

Spring & Summer 2023

Deep Learning with PyTorch

Tutorial and Lab Session Instructor (200-300 participants) at Global Mobile Internet Conference, Beijing, China

Spring 2018

Middle-School Mathematics and English

Volunteer Teaching for Low-Income Students in Northwestern China

Summer 2011

SERVICES

Reviewer ICML 2020-2024, NeurIPS 2020-2023, ICLR 2022, RLC 2024, CVPR 2021, TMLR, TPAMI, GCRL Workshop 2023.

Workshop Organizer Goal-Conditioned Reinforcement Learning (GCRL) Workshop at NeurIPS 2023.
The First Dataset Distillation Challenge at ECCV 2024.

OPEN-SOURCE PROJECTS_____ (106k stars on GitHub combined over projects that I made significant contributions to)

PyTorch 2017-2020

Core dev on data loading pipelines, CUDA/CPU kernels, ML ops, API design, autograd optimization, Python bindings, etc.

CycleGAN and pix2pix in PyTorch 2018-PRESENT

Maintainer for image-to-image translation toolkit

torchreparam 2019-2020

Built one of the first tools to reparametrize neural nets; now part of the large-scale SFDP training framework fairscale

torchqmet 2022-PRESENT

Built the first toolkit for learning quasimetrics with neural nets

HONORS AND AWARDS_____

Meta Ph.D. Fellowship Finalist 2023

Outstanding Reviewer for ICML 2022 2022

Top Reviewer for ICML 2020 2020

Merrill Lynch Graduate Fellowship 2019

UC Berkeley High Distinction in General Scholarship 2017

Best Summer Social Practice of Shanghai for my volunteer teaching in northwestern China 2011

OTHER ENGINEERING EXPERIENCES_____

Airbnb, Inc. 2016

Intern on Machine Learning Infrastructure

Facebook, Inc. 2015

Intern on Ads API Platform

Grue, Inc. 2015

Co-Founder

Sellegit, Inc. 2014

Intern on Mobile App Development

PUBLICATIONS (COMPLETE LIST)_____ (* indicates equal contribution)

A Framework for Evaluating the Efficacy of Foundation Embedding Models in Healthcare

Sonnet Xu, Haiwen Gui, Veronica Rotemberg, **Tongzhou Wang**, Yiqun Chen, Roxana Daneshjoui 2024

• [medRxiv](#)

The Platonic Representation Hypothesis

Minyoung Huh*, Brian Cheung*, **Tongzhou Wang***, Phillip Isola* 2024

• International Conference on Machine Learning 2024 [**ICML 2024 (Position Paper)**]

• [Code](#) [Webpage](#) [arXiv](#)

Optimal Goal-Reaching Reinforcement Learning via Quasimetric Learning

Tongzhou Wang, Antonio Torralba, Phillip Isola, Amy Zhang 2023

• International Conference on Machine Learning 2023 [**ICML 2023**].

• [Code](#) [Webpage](#) [arXiv](#)

Generalizing Dataset Distillation via Deep Generative Prior

George Cazenavette, **Tongzhou Wang**, Antonio Torralba, Alexei A. Efros, Jun-Yan Zhu 2023

• Conference on Computer Vision and Pattern Recognition 2023 [**CVPR 2023**].

• [Code](#) [Webpage](#) [arXiv](#)

Deep Augmentation: Enhancing Self-Supervised Learning through Transformations in Higher Activation Space

Rickard Br  l-Gabrielsson, [Tongzhou Wang](#), Manel Baradad, Justin Solomon

2023

- [arXiv](#)

Steerable Equivariant Representation Learning

Sangnie Bhardwaj, Willie McClinton, [Tongzhou Wang](#), Guillaume Lajoie, Chen Sun, Phillip Isola, Dilip Krishnan

2023

- [arXiv](#)

Improved Representation of Asymmetrical Distances with Interval Quasimetric Embeddings

[Tongzhou Wang](#), Phillip Isola

2022

- Workshop on Symmetry and Geometry in Neural Representations at NeurIPS 2022 [[NeurReps Workshop at NeurIPS 2022](#)].
- [PyTorch Package for Quasimetric Learning](#) [Webpage](#) [OpenReview](#) [arXiv](#)

Procedural Image Programs for Representation Learning

Manel Baradad, Chun-Fu Chen, Jonas Wulff, [Tongzhou Wang](#), Rogerio Feris, Antonio Torralba, Phillip Isola

2022

- Advances in Neural Information Processing Systems 2022 [[NeurIPS 2022](#)].
- [Code & Datasets](#) [Webpage](#) [OpenReview](#) [arXiv](#)

Denoised MDPs: Learning World Models Better Than the World Itself

[Tongzhou Wang](#), Simon S. Du, Antonio Torralba, Phillip Isola, Amy Zhang, Yuandong Tian

2022

- International Conference on Machine Learning 2022 [[ICML 2022](#)].
- [Code](#) [Webpage](#) [arXiv](#)

On the Learning and Learnability of Quasimetrics

[Tongzhou Wang](#), Phillip Isola

2022

- International Conference on Learning Representations 2022 [[ICLR 2022](#)].
- [Code](#) [Webpage](#) [OpenReview](#) [arXiv](#)

Dataset Distillation by Matching Training Trajectories

George Cazenavette, [Tongzhou Wang](#), Antonio Torralba, Alexei A. Efros, Jun-Yan Zhu

2022

- Conference on Computer Vision and Pattern Recognition 2022 [[CVPR 2022](#)].
- [Code](#) [Webpage](#) [arXiv](#)

Wearable ImageNet: Synthesizing Tileable Textures via Dataset Distillation

George Cazenavette, [Tongzhou Wang](#), Antonio Torralba, Alexei A. Efros, Jun-Yan Zhu

2022

- 5th Workshop on Computer Vision for Fashion, Art, and Design at CVPR 2022 [[CVFAD Workshop at CVPR 2022](#)].
- [Code](#) [Webpage](#) [Paper](#)

Totems: Physical Objects for Verifying Visual Integrity

Jingwei Ma, Lucy Chai, Minyoung Huh, [Tongzhou Wang](#), Ser-Nam Lim, Phillip Isola, Antonio Torralba

2022

- European Conference on Computer Vision 2022 [[ECCV 2022](#)].
- [Code](#) [Webpage](#) [arXiv](#)

Learning to See by Looking at Noise

Manel Baradad*, Jonas Wulff*, [Tongzhou Wang](#), Phillip Isola, Antonio Torralba

2021

- Advances in Neural Information Processing Systems 2021 [[NeurIPS 2021](#)].
- [Code & Datasets](#) [Webpage](#) [arXiv](#)

Understanding Contrastive Representation Learning through Alignment and Uniformity on the Hypersphere

[Tongzhou Wang](#), Phillip Isola

2020

- International Conference on Machine Learning 2020 [[ICML 2020](#)].
- [Code](#) [Webpage](#) [arXiv](#)

Rewriting a Deep Generative Model

David Bau, Steven Liu, [Tongzhou Wang](#), Jun-Yan Zhu, Antonio Torralba

2020

- European Conference on Computer Vision 2020 [[ECCV 2020](#)].
- [Code](#) [Webpage](#) [arXiv](#)

Diverse Image Generation via Self-Conditioned GANs

Steven Liu, [Tongzhou Wang](#), David Bau, Jun-Yan Zhu, Antonio Torralba

2020

- Conference on Computer Vision and Pattern Recognition 2020 [[CVPR 2020](#)].
- [Code](#) [Webpage](#) [arXiv](#)

Dataset Distillation

Tongzhou Wang, Jun-Yan Zhu, Antonio Torralba, Alexei A. Efros

2018

- [↗ Code](#) [Webpage](#) [arXiv](#)

Meta-Learning MCMC Proposals

Tongzhou Wang, Yi Wu, David A. Moore, Stuart J. Russell

2017

- Advances in Neural Information Processing Systems 2018 [[NeurIPS 2018](#)].
- Automatic Machine Learning Workshop at ICML 2017 (Oral) [[AutoML Workshop at ICML 2017 \(Oral\)](#)].
- [↗ arXiv](#)

Learning to Synthesize a 4D RGBD Light Field from a Single Image

Pratul Srinivasan, Tongzhou Wang, Ashwin Sreelal, Ravi Ramamoorthi, Ren Ng

2017

- International Conference on Computer Vision 2017 [[ICCV 2017](#)].
- [↗ Code](#) [arXiv](#)